

Umeå University
Department of Computing Science

Parallel Programming 7.5 p
5DV152

Exercises, Chapter/Topic 1

Submitted 2017-01-26
Author: Lorenz Gerber (dv15lgr@cs.umu.se lozger03@student.umu.se)
Instructor: Lars Karlsson / Mikael Ränner

Contents

1	Introduction	1
2	1.1 - Formulas for block partitioning	1
3	1.2 - Modify 1.1 with non-uniform costs	1
	3.1 First attempt	1
	3.2 Second attempt	1
	3.3 Third attempt	2
	3.4 Fourth attempt	2
4	1.3 - Tree-structured global sum	2
5	1.4 - Alternative algorithm for 1.3	3
6	1.5 - Generalization of 1.3 and 1.4	4
	6.1 Generalized form of 1.3	4
	6.2 Generalized form of 1.4	4
7	1.6 - Cost analysis of global sum algorithms	5
	References	5

1 Introduction

This report is part of the mandatory coursework. It describes the solutions for several chosen exercises from the course book [?].

2 1.1 - Formulas for block partitioning

The overwhelming idea is to load balance p number of cores with n tasks. Here, we use two functions to obtain block partitioning using a `for` loop:

```
for (my_i = my_first_i; my_i < my_last_i; my_i++)
```

The functions `my_first_i` and `my_last_i` are used to set the limits in the loop. Besides n , i and p we also need an index for the actual core: k . It is understood that indices i and k start at 0. The book text hints to start with the case when n is evenly divisible by p :

```
my_first_i = k * n / p
my_last_i = (k + 1) * n / p
```

Testing this expression for $n = 10$, $p = 5$, $k = \{0, 1, \dots, 4\}$ seems to be correct. Now when n is not even divisible by p , one has to distribute the $n \bmod p$ tasks for example on the first $n \bmod p$ cores:

```
my_first_i = k * n / p + (k < n mod p ? k : n mod p)
my_last_i = (k + 1) * n / p + (k + 1 < n mod p ? k + 1 : n mod p)
```

Testing this expression for $n = 9$, $p = 5$, $k = \{0, 1, \dots, 4\}$ gives the correct results.

3 1.2 - Modify 1.1 with non-uniform costs

In my opinion, this question can be understood in several different ways:

3.1 First attempt

Let there be 10 tasks (n) that take 2, 4, ..., 20 ms to finish, and there are 3 cores (p). So one could try to distribute the tasks onto the cores that the standard deviation of total runtimes for all cores is as little as possible. A possible solution for the above described actual setting would be:

i_9 to k_0 , i_8 to k_1 , i_7 to k_2 . Then, i_0 to k_0 , i_1 to k_1 , i_2 to k_2 . And i_6 to k_0 , i_5 to k_1 , i_4 to k_2 . Finally, i_3 to k_0 . This somehow works, makes sense and I guess it would be possible to devise a general formula for this.

However, it does not adhere to the actual question: *How would you change your answer to the preceding question if* And the preceding question is: *Devise formulas for the functions that calculate 'my_first_i' and 'my_last_i' in the global sum example.* Hence, to make formulas for first i and last i the tasks should be assigned to the cores in a consecutive sequence. Therefore, second attempt:

3.2 Second attempt

If the tasks have to be in a consecutive order, I would calculate the total amount of time all tasks together will take. That should be $n^2 + n$. This can be divided by the number of

cores available $\frac{n^2+n}{p} = r$. Now I would write an algorithm that iterates through n and assigns the tasks to the cores until r is 'reached'. Probably, there would be some edge case handling needed and it could make sense to iterate from high to low i 's.

To come up with a non-iterative approach for `my_first_i` and `my_last_i` formulas will be more mathematically involved. I could think of an approach where one needs to solve partial integrals of total time elapsed.

However, this way of asking the question somehow does not make much sense: The solution will be less optimal than the one devised in the first idea. So I read again the question and came up with again another interpretation:

3.3 Third attempt

The question states $i = k$ requires $k + 1$ times as much time than call $i = 0$. Somehow this does not make sense. Because in this formula, to define the time, there need to be equal indices k as indices i . That would mean that the problem is by definition embarrassingly parallel.

From the wording of the question in the second part *the first call takes..., the second call requires...* etc., I conclude that the author means that every time the algorithm is executed it will take inherently (in the current case) 2 milliseconds longer to finish. Maybe it should not be k , as index for cores, but just an arbitrary index to construct a number series. This brings me to the fourth and last idea:

3.4 Fourth attempt

If every call to the algorithm takes inherently longer, the task's i are actually independent of the time it takes to run and one can not start executing the 'long running' i 's as described in attempt 1 and 2. In this case, it's obvious that the solution for 1.1 is also valid here.

4 1.3 - Tree-structured global sum

The aim was to write pseudo code that calculates the tree structured global sum described on page 5. The book hints to use a variable called `divisor` that is initialized with the value 2 and another variable called `core_difference` that is initialized with the value 1. It was proposed that `divisor` is doubled in each iteration and that $n \bmod \text{divisor}$ is used to determine `send = 0` and `receive = 1`. From figure 1.1 in the book, one can see that this rule works. Note especially that $\text{core } k = 0$ will read in each iteration as $0 \bmod x = 0$ is true for any x . Further, the book proposes that `core_diff`, when doubled in each iteration, can be used to describe the difference in value between a core pair that is involved in a 'send/receive' operation. The correctness of this can also be verified in figure 1.1 of the book.

Assuming that k is the core index, this allows already to write the main part of the algorithm:

```
if( k % divisor == 0 )
    receive and add from k + core_diff
else
    send to k - core_diff
    break
```

Further, there need to be some control structure to know when to finish. Here it was decided to use a while loop with the comparison `core_diff < p`, where `p` is the number of cores in the system. In each iteration `divisor` and `core_diff` are doubled. Moreover, before exit, the core `k = 0` should return the result. Below is the complete code listing that incorporates the the described features:

```
divisor = 2
core_diff = 1

while(core_diff < p){
    if( k % divisor == 0 )
        receive and add from k + core_diff
    else
        send to k - core_diff
        break

    divisor * 2
    core_diff * 2
}

if( k == 0 )
    return final result
```

5 1.4 - Alternative algorithm for 1.3

Here the aim was to modify the pseudo code from 1.3 to use bit shift operators, basically to obtain `k` indices for send and receive. In the book, the idea is visually described using a table (page 13, exercise 1.4). The bitwise `Xor` operator is applied to a binary bitmask with the initial value 001_2 and the binary value of each `k`. This will flip the last bit in every `k` value, hence always two `k` values exchange their initial value with each other. This can be exploited to use the core `k` whose value becomes lower by the bitshift as the sender and the other one in the flipping pair as receiver.

```
counter = 2
bitmask = 1

while ( counter < p ){
    if(k bitwiseXor bitmask > k)
        receive and add from k bitwiseXor bitmask
    else
        send to k bitwiseXor bitmask
        break

    bitwiseLeft bitmask
    counter * 2
}

if( k == 0)
```

4(5)

```
return final result
```

6 1.5 - Generalization of 1.3 and 1.4

6.1 Generalized form or 1.3

```
divisor = 2
core_diff = 1

while(core_diff < p){
    if(k % divisor == 0)
        if (k + core_diff < p)
            receive and add from k + core_diff
        else
            send to k - core_diff
            break

    divisor * 2
    core_diff * 2
}

if( k == 0 )
    return final result
```

6.2 Generalized form of 1.4

```
counter = 2
bitmask = 1

while ( counter < p ){
    if(k bitwiseXor bitmask > k)
        if (k bitwiseXor bitmask < p)
            receive and add from k bitwiseXor bitmask
        else
            send to k bitwiseXor bitmask
            break

    bitwiseLeft bitmask
    counter * 2
}

if( k == 0)
    return final result
```

7 1.6 - Cost analysis of global sum algorithms

The number of receives and additions for core 0 is in the original 'pseudo-code global sum' $p-1$ and in the tree-structured global sum $\text{ceiling}(\log_2(p))$.