

Umeå University
Department of Computing Science

Parallel Programming 7.5 p
5DV152

Exercises, Chapter/Topic 1

Submitted 2017-02-16
Author: Lorenz Gerber (dv15lgr@cs.umu.se lozger03@student.umu.se)
Instructor: Lars Karlsson / Mikael Ränner

Contents

1	Introduction	1
2	3.2 - Generalization of algorithm for trapezoidal rule	1
3	3.6 - Array distributions	2
	3.1 Block distribution	2
	3.2 Cyclic distribution	2
	3.3 Block cyclic distribution	2
4	3.8 - Tree-structured algorithms for scatter and gather	2
5	3.9 - Vector scaling and dot product	2
6	3.11 - Prefix sums	2
7	3.13 - Generalization of vector scaling and dot product	2
8	3.16 - Diagram for a butterfly implementation of allgather	2
9	3.18 - Derived data types	2
10	3.20 - Pack and unpack	2
11	3.21 - Matrix-vector multiplication	2
12	3.22 - Timing the trapezoidal rule	2
13	3.27 - Speedup and efficiency of odd-even sort	3
	References	3

1 Introduction

This report is part of the mandatory coursework. It describes the solutions for several chosen exercises from the course book [?].

2 3.2 - Generalization of algorithm for trapezoidal rule

Two functions to adapt the *trapezoidal rule* for `calc_local_a` and `calc_local_b` were written and tested with the source code from the book (*mpi_trap.c*).

```
double calc_local_a(int my_rank, double a, double b, int n, int comm_sz){
    double local_a = 0;
    double h = 0;
    int local_n = 0;
    int rest_n = 0;

    h = (b-a)/n;
    local_n = n/comm_sz;

    rest_n = n%comm_sz;

    if(my_rank < rest_n){
        local_a = a + my_rank*local_n*h + my_rank*h;
    } else {
        local_a = a + my_rank*local_n*h + rest_n*h;
        local_a += (my_rank-rest_n) * h;
    }

    return local_a;
}

double calc_local_b(int my_rank, double a, double b, int n, int comm_sz){
    double h;
    int local_n;

    h = (b-a)/n;
    local_n = n/comm_sz;

    if (my_rank == (comm_sz-1)){
        return a + my_rank+1*local_n*h;
    } else {
        return calc_local_a(my_rank+1, a, b, n, comm_sz);
    }
}
```

3 3.6 - Array distributions

Given is a vector x of length n with the indices i . The number of processes used is `comm_sz` and the index of the current process is `my_rank`.

3.1 Block distribution

For a block distribution, we devise the functions `local_start_i` and `local_end_i`.

```
local_start_i:
my_rank * n / comm_sz + (my_rank < n mod comm_sz ? my_rank : n mod comm_sz)

local_end_i:
(my_rank + 1) * n / comm_sz + (my_rank + 1 < n mod comm_sz ? my_rank + 1 : n mod comm_sz)
```

3.2 Cyclic distribution

3.3 Block cyclic distribution

4 3.8 - Tree-structured algorithms for scatter and gather

5 3.9 - Vector scaling and dot product

takes a while to solve, requires programming

6 3.11 - Prefix sums

takes a while to solve requires programming

7 3.13 - Generalization of vector scaling and dot product

8 3.16 - Diagram for a butterfly implementation of allgather

9 3.18 - Derived data types

takes a while to solve requires programming

10 3.20 - Pack and unpack

requires programming

11 3.21 - Matrix-vector multiplication

takes a while to solve requires programming requires testing

12 3.22 - Timing the trapezoidal rule

takes a while to solve Requires programming requires testing

13 3.27 - Speedup and efficiency of odd-even sort