

**Umeå University**  
Department of Computing Science

**Parallel Programming 7.5 p**  
**5DV152**

**Exercises, Chapter/Topic 3**

Submitted 2017-02-16  
Author: Lorenz Gerber (dv15lgr@cs.umu.se lozger03@student.umu.se)  
Instructor: Lars Karlsson / Mikael Ränner

## Contents

1	Introduction	1
2	3.2 - Generalization of algorithm for trapezoidal rule	1
3	3.6 - Array distributions	1
4	3.8 - Tree-structured algorithms for scatter and gather	1
5	3.9 - Vector scaling and dot product	1
6	3.11 - Prefix sums	1
7	3.13 - Generalization of vector scaling and dot product	6
8	3.16 - Diagram for a butterfly implementation of allgather	6
9	3.18 - Derived data types	6
10	3.20 - Pack and unpack	7
11	3.21 - Matrix-vector multiplication	7
	11.1 Speed comparison	7
	11.2 Observed variabilities	7
	11.3 Cluster around Minmum, mean, median	7
12	3.22 - Timing the trapezoidal rule	7
13	3.27 - Speedup and efficienci of odd-even sort	11
	References	12
A	C Source Code for Exercise 3.2	12
B	C Source Code for Exercise 3.9	13
C	C Source Code for Exercise 3.11	15
D	C Source Code for Exercise 3.13	16
E	C Source code for Exercise 3.18	19
F	C Source Code for Exercise 3.20	21

## 1 Introduction

This report is part of the mandatory coursework. It describes the solutions for several chosen exercises from the course book [3].

## 2 3.2 - Generalization of algorithm for trapezoidal rule

Two functions to adapt the *trapezoidal rule* for `calc\_local\_a` and `calc\_local\_b` were written and tested with the source code from the book (*mpi\_trap.c*). The source can be found in appendix A.

## 3 3.6 - Array distributions

### Block distribution

Block distribution can be obtained by  $b = \lfloor i \div p \rfloor$  where  $b$  is the block number,  $i$  the index of  $n$  and  $p$  is the number of processes. This solution is however not fair. An improved, fair expression can be devised using ternary operators:

$$i < n \bmod p \times \lceil n \div p \rceil ? \lfloor i \div \lceil n \div p \rceil \rfloor : n \bmod p + \lfloor (i - n \bmod p \times \lceil n \div p \rceil) \div \lceil n \bmod p \rceil \rfloor$$

### Cyclic distribution

Cyclic distribution is described by  $b = i \bmod p$  with  $b$  as block number  $i$  as index of  $n$  and  $p$  as number of processes.

### Block cyclic distribution

Block cyclic distribution can be expressed as  $b = \lfloor i \div l \rfloor \bmod p$  where  $b$  is block index,  $i$  index of  $n$ ,  $l$  block length and  $p$  number of processes.

## 4 3.8 - Tree-structured algorithms for scatter and gather

The diagram for tree-structured scatter is shown in figure 1. The arrows show the communication events which for the present case of `comm_sz = 8` and `n = 16` is 7. This is also true for the tree-structured gather shown in figure 2.

## 5 3.9 - Vector scaling and dot product

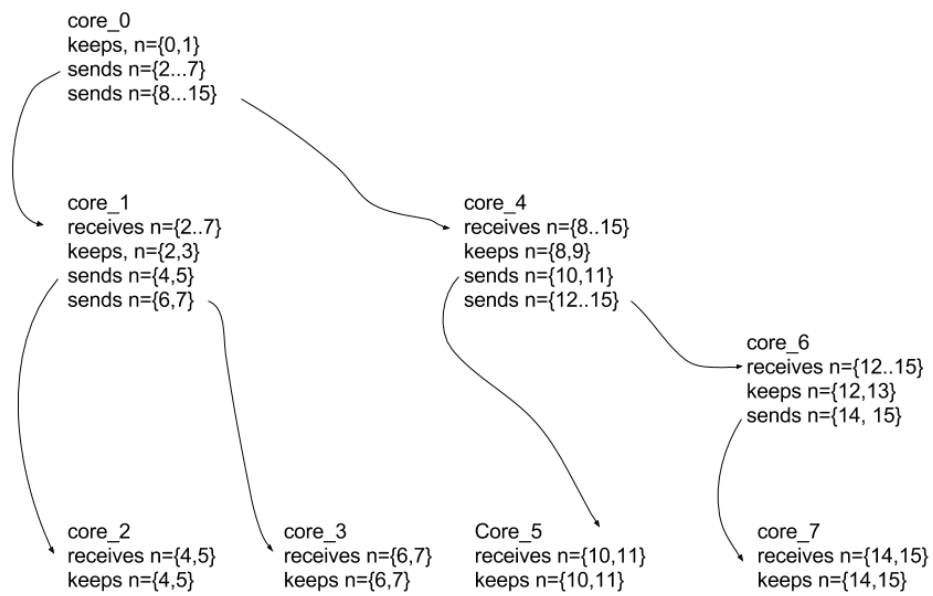
The Source code for the vector scaling and dot product MPI program can be found in appendix B.

## 6 3.11 - Prefix sums

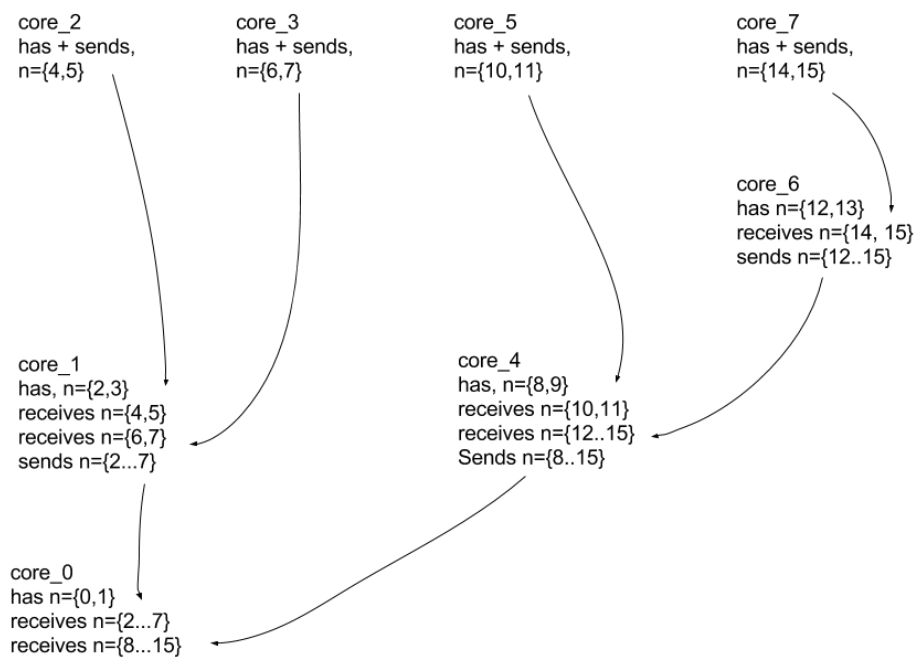
### Serial Algorithm

Assume vector  $x$  of length  $n$  and prefix sums vector  $y$  to be calculated.

```
y_0 = x_0
for (i = 1, i < n, i++)
```



**Figure 1:** This graph shows a tree based implementation of scatter for  $\text{comm\_sz} = 8$  and  $n = 16$ . It has 7 communication events.



**Figure 2:** This graph shows a tree based implementation of gather for  $comm\_sz = 8$  and  $n = 16$ . It has 7 communication events.

$$y_i = y_{i-1} + n_i$$

### Parallel Algorithm

After Blelloch [1], parallel prefix-sum, or ‘Sum Scan Algorithm’ can be calculated in two steps, the ‘Up-sweep’ (Reduce) and ‘Down-sweep’. Below, pseudo code from Valdez [5] where  $x$  is the input data,  $n$  the size of the input and  $d$  the number of processeors. If  $n \neq 2^k$ , then  $n$  has to be extended with *zero*’s. The ‘Up-Sweep’:

```
for d = 0 to log2(n) - 1 do
  for all k = 0 to n - 1 by 2^(d+1) in parallel do
    x[k + 2^(d+1) - 1] = x[k + 2^d - 1] + x[k + 2^(d+1) - 1]
```

The ‘Down-Sweep’:

```
x[n - 1] = 0
for d = log2(n) - 1 down to 0 do
  for all k = 0 to n - 1 by 2^(d+1) in parallel do
    t = x[k + 2^d - 1]
    x[k + 2^d - 1] = x[k + 2^(d+1) - 1]
    x[k + 2^(d+1) - 1] = t + x[k + 2^(d+1) - 1]
```

### Minimum Communication


An algorithm that uses for the prefix-sum of length  $n = 2^k$  only  $k$  communication steps can be derived from the communication model in a hypercube as shown in figure 3.

Below follows pseudo code for parallel hypercube prefix sum (from [2]):


```
procedure PREFIX_SUMS_HCUBE(my_id, my number, d, result)
begin
  result := my_number;
  msg := result;
  for i := 0 to d - 1 do
    partner := my_id XOR 2i;
    send msg to partner;
    receive number from partner;
    msg := msg + number;
    if (partner < my_id) then result := result + number;
  endfor;
end PREFIX_SUMS_HCUBE
```

### MPI Program using MPI\_Scan


A minimal working example using MPI\_Scan was implemented. The source code can be found in appendix C. An array of random numbers are generated on each proces. Then the prefix sums are collected onto process zero. The prefix sums of the arrays are collected index wise. Arrays of length  $n$  and  $p$  processes will result in a  $n \times p$  prefix sum array.




(a) Initial distribution of values



(b) Distribution of sums before second step




(c) Distribution of sums before third step



(d) Final distribution of prefix sums

**Figure 3:** This figure shows the calculation of prefix sums with minimal communication phases. The figure is take from [2, chap 4.3]



**Figure 4:** This figure shows a butterfly allgather implementation for 8 cores and a vector of length 16.

### 7 3.13 - Generalization of vector scaling and dot product

An MPI program for ‘Vector Scaling and Dot Product’ using ‘MPI\_Scatterv’/‘MPI\_Gatherv’ was written according to the specifications. The source code can be found in appendix D.

### 8 3.16 - Diagram for a butterfly implementation of allgather

Figure 4 shows an diagram for an *allgather* implementation using *butterfly* communication. The examples shows the communication steps for 8 processes and a data vector length of 16.

### 9 3.18 - Derived data types

Here `MPI_Type_vector` was used to implement a derived data type to represent a block cyclic data distribution. The current implementation assumes that the data is even divisible by `blocksize × processes`. This allows to use only one data type definition for the full length block cyclic structe of one process. The source code can be found in appendix E



**Table 1** Median runtimes for the Matrix-Vector Multiplication. The program was run on the HPC2N ‘abisko’ cluster. Processes were assigned to individual nodes.

comm_sz	Order of Matrix (milliseconds)				
	1024	2048	4096	8192	16384
1	0.01	0.02	0.03	0.07	0.14
2	0.01	0.02	0.02	0.04	0.08
4	0.01	0.02	0.02	0.03	0.04
8	0.02	0.02	0.02	0.02	0.03
16	0.03	0.03	0.03	0.03	0.05

## 10 3.20 - Pack and unpack

The MPI functions MPI\_Pack and MPI\_Unpack were used transfer complex datastructures. As example, a new *Get\_input* function for the ‘trapezoidal sum’ program was written that transfers the data from process 0 to all processes. The source code of the modified *Get\_input* function can be found in appendix F.

## 11 3.21 - Matrix-vector multiplication

The ‘Matrix-vector multiplication’ program from the course book was run on the HPC2N Abisko cluster for benchmarking and comparison to the values in the course book. Herefore the source *mpi\_mat\_vect\_time.c* provided from the coursebooks homepage was compiled with MPI-gcc compiler on the Abisko cluster. Problemsize and process number were chosen according to the example in the course book.

### 11.1 Speed comparison

The absolute run-times values differ significantly from the benchmarked system in the book, see [3, p. 123] and table 1. The bechmarking was run on the ‘abisko’ cluster with the setting to run one process per node. For such a small program, this was performance wise probably not the most optimal choice: It can be seen in table 1 that all runtimes for 16 processes are longer than those for respective problem size and fewer processes. Most likely, this is due to the higher communication overhead.

### 11.2 Observed variabilities

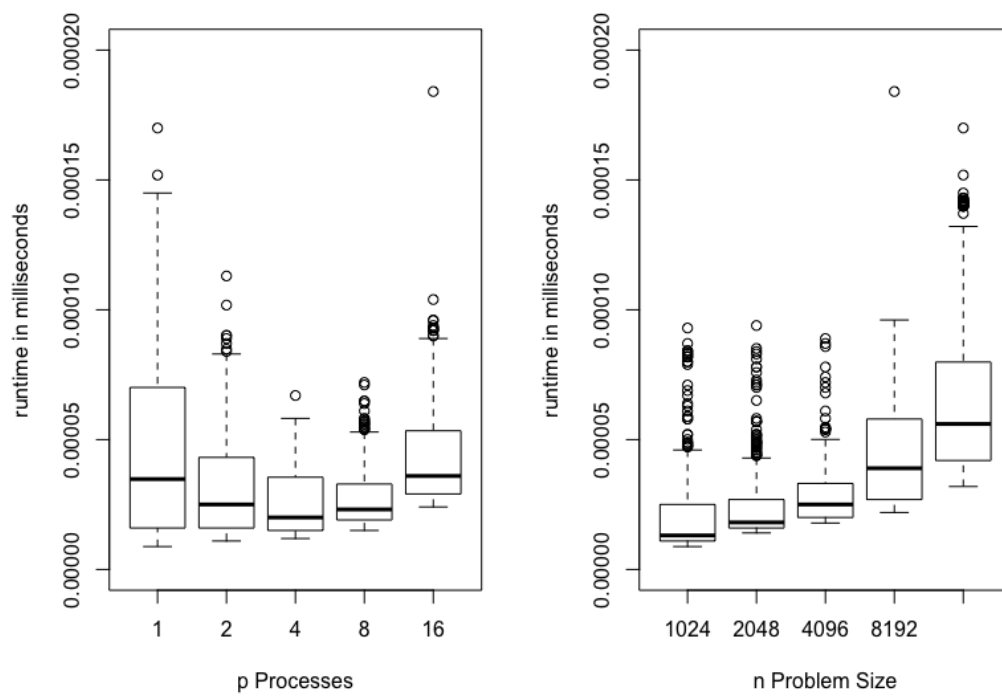
Variability was assessed using boxplots, see figure 5.

### 11.3 Cluster around Minmum, mean, median


For the clustering of mininum, mean and median, plot showing the timings normalized to their respective largest value were plotted. In each plot series, mean and median were also shown (Figure 6).

## 12 3.22 - Timing the trapezoidal rule

The problem size should be select in a way that the runtimes occur in a regime that is measurable with the given precision of the timing functionality. Further, it is favourable to have a series of multiples i.e the double, the quadruple etc. Here it was chose to use



**Figure 5:** This figure shows boxplot of runtimes in relation to number of processes and problem size.



**Figure 6:** This figure shows the normalized ranges of runtimes to compare minima, maxima, median and mean values. Mean values are marked by a red square and median values by a blue circle

**Table 2** This table shows the mean of the benchmarking values for the trapezoidal rule

processes	number of trapezoids				
	65536	131072	262144	524288	1048576
1	0.73	1.36	2.70	5.39	10.77
2	0.31	0.62	1.23	2.46	4.91
4	0.16	0.33	0.67	1.29	2.59
8	0.09	0.17	0.33	0.66	1.32
16	0.27	0.10	0.22	0.35	0.73

**Table 3** This table shows the median of the benchmarking values for the trapezoidal rule.

processes	number of trapezoids				
	65536	131072	262144	524288	1048576
1	0.67	1.34	2.70	5.38	10.77
2	0.31	0.61	1.22	2.46	4.91
4	0.17	0.33	0.67	1.32	2.59
8	0.09	0.17	0.34	0.67	1.33
16	0.06	0.10	0.18	0.35	0.70

the series 65536, 131072, 262144, 524288, 1048576 as problem size  $n$  for the number of trapezoids.

Here, the parameters for the batch runs on HPC2N's 'abisko' were chose to assign one processor for each process instead of a whole node for each process. This had certainly a positive effect on lowering the communication overhead for higher number of processes.

In most cases, median (table 3) and mean values (2) were just a few digits apart, with the median always being the lower. However, in some cases the sensitivity of the mean towards outliers can be see. For example for the mean time of  $p = 16$  and  $n = 65536$  where the mean is significantly off from the media value.

The minima values 4 are all consistently smaller than the median values 3, however it's not a large difference.

How does minima time compares to median.

The benchmark of the trapezoidal rule algorithm resulted for low number of process values in improbable speedup and efficiency values (too good). However, the serial algorithm was run on a normal machine that could be slower than the individual dedicated HPC2N machines. At a intermediate number of processes the algorithm has an almost linear speedup and respective efficiencies. At larger numbers of processes, the communication overhead degrades the performance a bit. However, the tested larger problem sizes could still make up

**Table 4** This table shows the minima values for the trapezoidal rule algorithm.

processes	number of trapezoids				
	65536	131072	262144	524288	1048576
1	0.67	1.33	2.62	5.34	10.72
2	0.29	0.58	1.16	2.33	4.86
4	0.15	0.30	0.60	1.19	2.51
8	0.08	0.15	0.30	0.60	1.20
16	0.05	0.09	0.18	0.34	0.66

**Table 5** This table shows the speedups for the trapezoidal rule algorithm.

processes	number of trapezoids				
	65536	131072	262144	524288	1048576
2	2.16	2.20	2.21	2.19	2.19
4	3.94	4.06	4.03	4.08	4.16
8	7.44	7.88	7.94	8.03	8.10
16	11.17	13.40	15.00	15.37	15.39

**Table 6** This table shows the efficiencies for the trapezoidal rule algorithm.

processes	number of trapezoids				
	65536	131072	262144	524288	1048576
2	1.08	1.10	1.11	1.09	1.10
4	0.99	1.02	1.01	1.02	1.04
8	0.93	0.99	0.99	1.00	1.01
16	0.70	0.84	0.94	0.96	0.96

for it. The algorithm is certainly weakly scalable, in some domains even strongly scalable.

### 13 3.27 - Speedup and efficiency of odd-even sort

First it should be noted that the comparison between serial and parallel timings is based on different algorithms. For serial timings, ‘quick sort’ was used and for the parallel ‘odd-even’ sort. Both algorithms have a worst case performance of  $O(n^2)$  and a best-case performance of  $O(n)$ . Average performance of quick sort is  $O(n \log n)$  while odd-even sort, when ignoring the communication overhead, should effectively perform at  $O(n)$  [4].

This observation fits well with the calculated speedups 7 and efficiencies 8: At low number of processes, the parallel implementation exceeds even the theoretical ideal of linear speedups (2 processes, speedups: 2.05, 2.09, 2.05, 2.02, 2.09) and correspondingly also the efficiencies (1.02 - 1.05). However, with increasing number of processes, the overhead for communication increases hence, both speedups and efficiencies degrade somewhat. From a theoretical point of view, the communication overhead at a certain number of processes is a constant, hence it should ‘fall away’ in the big O notation for big enough problem size. It can be seen that both speedups and efficiencies increase again at high number of processes for the largest problem sizes.

Hence, for low number of processes odd-even sort is strongly scalable. For larger numbers of processes, the performance degrades somewhat but weakly scalable sounds here

**Table 7** Speedups

processes	number of keys				
	200	400	800	1600	3200
2	2.05	2.09	2.05	2.02	2.09
4	4.00	4.13	4.06	4.15	4.19
8	7.33	7.92	7.65	7.55	8.18
16	11.73	13.57	13.45	13.83	13.85

**Table 8** *Efficiencies*

processes	number of keys				
	200	400	800	1600	3200
2	1.02	1.04	1.03	1.03	1.05
4	1.00	1.03	1.02	1.04	1.05
8	0.92	0.99	0.96	0.94	1.02
16	0.73	0.85	0.84	0.86	0.87

almost as an underestimation. However, by definition, it is weakly scalable.

## References

- [1] G.E. Blelloch. *Vector Models for Data-Parallel Computing*. MIT Press, Cambridge MA, USA, 1990.
- [2] A Grama, A Gupta, G Karypis, and V Kumar. *Introduction to Parallel Computing*. Pearson, Essex, England, 2003.
- [3] P.S. Pacheco. *An Introduction to Parallel Programming*. Morgan Kaufman, 2011.
- [4] Parallel Sorting, lecture notes kth, sf2568. <https://www.math.kth.se/na/SF2568/parpro-16/F7.pdf>. accessed: 2017-02-15.
- [5] Stackoverflow, parallel prefix sum - fastest implementation. <http://stackoverflow.com/questions/10053629/parallel-prefix-sum-fastest-implementation>, 2012. accessed: 2017-02-04.

## A C Source Code for Exercise 3.2

```
double calc_local_a(int my_rank, double a, double b, int n, int comm_sz){
    double local_a = 0;
    double h = 0;
    int local_n = 0;
    int rest_n = 0;

    h = (b-a)/n;
    local_n = n/comm_sz;

    rest_n = n%comm_sz;

    if(my_rank < rest_n){
        local_a = a + my_rank*local_n*h + my_rank*h;
    } else {
        local_a = a + my_rank*local_n*h + rest_n*h;
        local_a += (my_rank-rest_n) * h;
    }
}
```

```

    return local_a;
}

double calc_local_b(int my_rank, double a, double b, int n, int comm_sz){
    double h;
    int local_n;

    h = (b-a)/n;
    local_n = n/comm_sz;

    if (my_rank == (comm_sz-1)){
        return a + my_rank+1*local_n*h;
    } else {
        return calc_local_a(my_rank+1, a, b, n, comm_sz);
    }
}

```

## B C Source Code for Exercise 3.9

```

#include <stdio.h>
#include <mpi.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    int my_rank, comm_sz;
    int n, local_n, local_dotp_sum = 0, scalar, result_dot;
    int* local_vec1;
    int* local_vec2;
    int* vector1;
    int* vector2;

    /* Initializing */
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    srand(time(NULL));

    /* Obtaining Data */
    if(my_rank==0 && argc > 1){

        if(strcmp(argv[1], "r") == 0){

```

```

    printf("using random data, vector length = %d\n", 100*comm_sz);
    n = 100*comm_sz;
    vector1 = (int *) malloc(100*comm_sz * sizeof(int));
    vector2 = (int *) malloc(100*comm_sz * sizeof(int));

    for(int i = 0; i < n; i++){
vector1[i] = rand() % 1000;
vector2[i] = rand() % 1000;
    }
    scalar = rand() % 1000;
}

} else if (my_rank==0){

    printf("enter vector length\n");
    scanf("%d", &n);
    printf("enter integer vector 1\n");

    vector1 = (int *) malloc(n * sizeof(int));
    vector2 = (int *) malloc(n * sizeof(int));

    for(int i = 0; i < n; i++){
        scanf("%d", &vector1[i]);
    }

    printf("enter integer vector 2\n");

    for(int i = 0; i < n; i++){
        scanf("%d", &vector2[i]);
    }

    printf("enter integer scalar\n");
    scanf("%d", &scalar);
}

/* Distribute Data */
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

local_n = n/comm_sz;
local_vec1 = (int*) malloc(local_n * sizeof(int));
local_vec2 = (int*) malloc(local_n * sizeof(int));

MPI_Bcast(&scalar, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Scatter(vector1, local_n, MPI_INT, local_vec1, local_n, MPI_INT, 0, MPI_COMM_WORLD)
MPI_Scatter(vector2, local_n, MPI_INT, local_vec2, local_n, MPI_INT, 0, MPI_COMM_WORLD)

/* Calculations */
/* Calculate Dot Product */

```



```

for(int i = 0; i< local_n; i++){
    local_vec2[i]*=local_vec1[i];
}

/* Calculate vector-scalar product */
for(int i = 0; i< local_n; i++){
    local_vec1[i]*=scalar;
}

/* Summing for dot product */
for(int i = 0; i < local_n; i++){
    local_dotp_sum += local_vec2[i];
}

/* Collect Data */
MPI_Gather(local_vec1, local_n, MPI_INT, vector1, local_n, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Reduce(&local_dotp_sum, &result_dot, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

/* Results */
if(my_rank == 0){
    printf("dot product = %d\n", result_dot);

    printf("vector-scalar product = ");
    for(int i = 0; i < n;i++){
        printf("%d ", vector1[i]);
    }
    printf("\n");
}

/* Clean up */
if(my_rank==0){
    free(vector1);
    free(vector2);
}

free(local_vec1);
free(local_vec2);

MPI_Finalize();

return 0;
}

```

## C C Source Code for Exercise 3.11

```

#include <stdio.h>
#include <stdlib.h>

```

16(22)

```
#include <mpi.h>
#include <time.h>

int main(void){

    int my_rank, comm_sz;
    int* initial_vector;
    int* prefix_sums;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
    srand(time(NULL));

    initial_vector = (int *) malloc(10*sizeof(int));
    prefix_sums = (int *) calloc(10, sizeof(int));

    for(int i = 0; i < 10;i++){
        initial_vector[i] = rand() % 1000;
    }

    MPI_Scan(initial_vector, prefix_sums, 10,MPI_INT, MPI_SUM, MPI_COMM_WORLD);

    for(int i = 0; i< comm_sz;i++){
        if (i == my_rank){
            for(int j = 0;j < 10;j++){
                printf("%d ", prefix_sums[j]);
            }
        }
    }
    printf("\n");

    MPI_Finalize();
    return 0;
}
```

## **D C Source Code for Exercise 3.13**

```
#include <stdio.h>
#include <mpi.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
```

```

int my_rank, comm_sz;
int n, local_n, local_dotp_sum = 0, scalar, result_dot;
int* sendcounts;
int* displs;
int* local_vec1;
int* local_vec2;
int* vector1;
int* vector2;

/* Initializing */
MPI_Init(NULL, NULL);
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);
srand(time(NULL));

/* Obtaining Data */
if(my_rank==0 && argc > 1){

    if(strcmp(argv[1], "r") == 0){
        printf("using random data, vector length = %d\n", 100*comm_sz);
        n = 100*comm_sz;
        vector1 = (int *) malloc(100*comm_sz * sizeof(int));
        vector2 = (int *) malloc(100*comm_sz * sizeof(int));

        for(int i = 0; i < n; i++){
            vector1[i] = rand() % 1000;
            vector2[i] = rand() % 1000;
        }
        scalar = rand() % 1000;
    }

    } else if (my_rank==0){

        printf("enter vector length\n");
        scanf("%d", &n);
        printf("enter integer vector 1\n");

        vector1 = (int *) malloc(n * sizeof(int));
        vector2 = (int *) malloc(n * sizeof(int));

        for(int i = 0; i < n; i++){
            scanf("%d", &vector1[i]);
        }

        printf("enter integer vector 2\n");

        for(int i = 0; i < n; i++){
            scanf("%d", &vector2[i]);

```

```

    }

    printf("enter integer scalar\n");
    scanf("%d", &scalar);
}

/* Distribute Data */
MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);

/* Fixing sendcounts for general n */
sendcounts = (int *) malloc(comm_sz * sizeof(int));
displs = (int *) calloc(comm_sz, sizeof(int));

for(int i = 0; i < comm_sz; i++){
    if(n % comm_sz > i){
        sendcounts[i] = (n/comm_sz) + 1;
    } else {
        sendcounts[i] = (n/comm_sz);
    }
}

for(int i = 1; i < comm_sz; i++){
    displs[i] = displs[i-1] + sendcounts[i];
}

local_n = sendcounts[my_rank];

local_vec1 = (int*) malloc(local_n * sizeof(int));
local_vec2 = (int*) malloc(local_n * sizeof(int));

MPI_Bcast(&scalar, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Scatterv(vector1, sendcounts, displs, MPI_INT, local_vec1, sendcounts[my_rank], MPI_COMM_WORLD);
MPI_Scatterv(vector2, sendcounts, displs, MPI_INT, local_vec2, sendcounts[my_rank], MPI_COMM_WORLD);

/* Calculations */
/* Calculate Dot Product */
for(int i = 0; i < local_n; i++){
    local_vec2[i] *= local_vec1[i];
}

/* Calculate vector-scalar product */
for(int i = 0; i < local_n; i++){
    local_vec1[i] *= scalar;
}

```

```

/* Summing for dot product */
for(int i = 0; i < local_n; i++){
    local_dotp_sum += local_vec2[i];
}

/* Collect Data */
MPI_Gatherv(local_vec1, sendcounts[my_rank], MPI_INT, vector1, sendcounts, displs, MPI_
MPI_Reduce(&local_dotp_sum, &result_dot, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

/* Results */
if(my_rank == 0){
    printf("dot product = %d\n", result_dot);

    printf("vector-scalar product = ");
    for(int i = 0; i < n;i++){
        printf("%d ", vector1[i]);
    }
    printf("\n");
}

/* Clean up */
if(my_rank==0){
    free(vector1);
    free(vector2);
}

free(local_vec1);
free(local_vec2);

MPI_Finalize();

return 0;
}

```

## E C Source code for Exercise 3.18

```

#include <stdio.h>
#include <mpi.h>
#include <stdlib.h>

```

```

int Read_vector(int block_length, double* vector, double** slice, int vector_length, int
int Print_vector(int block_length, double** slice, double* vector, int vector_length, int
int main(void) {

```

```

int my_rank, comm_sz;
double vec[18] = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17};
double readback[18] = { 0 };
double * slice;

MPI_Init(NULL, NULL);
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &comm_sz);

Read_vector(2, vec, &slice, 18, my_rank, comm_sz);
Print_vector(2, &slice, readback, 18, my_rank, comm_sz);

/* Print out the result */
if(my_rank == 0){
    printf("The vector after round Read_vector and Print_vector on Core 0:\n");
    for(int i = 0; i < 18; i++){
        printf("i %d: %f\n", i, readback[i]);
    }
}

/* Clean up */
MPI_Finalize();
free(slice);
return 0;
}

int Read_vector(int block_length, double* vector, double** slice, int vector_length, int my_rank, int comm_sz)
{
    /* calc block cyclic distro */
    int no_cycles = vector_length / (block_length * comm_sz);
    int cycle_block_length = block_length * comm_sz;
    int slice_length = no_cycles * block_length;

    /* define data type */
    MPI_Datatype block_cyclic;
    MPI_Type_vector(no_cycles, block_length, cycle_block_length, MPI_DOUBLE, &block_cyclic);
    MPI_Type_commit(&block_cyclic);
    MPI_Status info;

    if (my_rank == 0){
        for (int i = 1; i < comm_sz; i++){
            MPI_Send(&vector[i*block_length], 1, block_cyclic, i, 99, MPI_COMM_WORLD);
        }
        *slice = (double *) calloc(slice_length, sizeof(double));
        MPI_Sendrecv(&vector[0], 1, block_cyclic, 0, 99,
            *slice, slice_length, MPI_DOUBLE, 0, 99, MPI_COMM_WORLD, &info);
    }
}

```

```

    } else {
        *slice = (double *) calloc(slice_length, sizeof(double));
        MPI_Recv(*slice, slice_length, MPI_DOUBLE, 0, 99, MPI_COMM_WORLD, &info);
    }

    /* Clean up */
    MPI_Type_free(&block_cyclic);

    return 0;
}

int Print_vector(int block_length, double** slice, double* vector, int vector_length, int

    /* calc block cyclic distro */
    int no_cycles = vector_length / (block_length * comm_sz);
    int cycle_block_length = block_length * comm_sz;
    int slice_length = no_cycles * block_length;

    /* define data type */
    MPI_Datatype block_cyclic;
    MPI_Type_vector(no_cycles, block_length, cycle_block_length, MPI_DOUBLE, &block_cyclic)
    MPI_Type_commit(&block_cyclic);
    MPI_Status info;

    if(my_rank != 0){
        MPI_Send(*slice, slice_length, MPI_DOUBLE, 0, 99, MPI_COMM_WORLD);

    } else {
        for (int i = 1; i < comm_sz; i++){
            MPI_Recv(&vector[i*block_length], 1, block_cyclic, i, 99, MPI_COMM_WORLD, &info);
        }
        MPI_Sendrecv(*slice, slice_length, MPI_DOUBLE, 0, 99, &vector[0], 1,
block_cyclic, 0, 99, MPI_COMM_WORLD, &info);
    }

    /* Clean up */
    MPI_Type_free(&block_cyclic);

    return 0;
}

```

## F C Source Code for Exercise 3.20

```

/*-----
* Function:      Get_input
* Purpose:       Get the user input:  the left and right endpoints

```

22(22)

```
*          and the number of trapezoids
* Input args:  my_rank:  process rank in MPI_COMM_WORLD
*              comm_sz:  number of processes in MPI_COMM_WORLD
* Output args: a_p:  pointer to left endpoint
*              b_p:  pointer to right endpoint
*              n_p:  pointer to number of trapezoids
*/
void Get_input(int my_rank, int comm_sz, double* a_p, double* b_p,
               int* n_p) {

    char pack_buf[100];
    int position = 0;

    if (my_rank == 0) {
        printf("Enter a, b, and n\n");
        scanf("%lf %lf %d", a_p, b_p, n_p);

        MPI_Pack(&a_p, 1, MPI_DOUBLE, pack_buf, 100, &position, MPI_COMM_WORLD);
        MPI_Pack(&b_p, 1, MPI_DOUBLE, pack_buf, 100, &position, MPI_COMM_WORLD);
        MPI_Pack(&n_p, 1, MPI_INT, pack_buf, 100, &position, MPI_COMM_WORLD);
    }

    MPI_Bcast(pack_buf, 100, MPI_PACKED, 0, MPI_COMM_WORLD);

    if(my_rank != 0) {
        MPI_Unpack(pack_buf, 100, &position, a_p, 1, MPI_DOUBLE, MPI_COMM_WORLD);
        MPI_Unpack(pack_buf, 100, &position, b_p, 1, MPI_DOUBLE, MPI_COMM_WORLD);
        MPI_Unpack(pack_buf, 100, &position, n_p, 1, MPI_INT, MPI_COMM_WORLD);
    }
} /* Get_input */
```