Umeå University
Department of Computing Science

Parallel Programming 7.5 p
5DV152

Exercises, Chapter/Topic 4

Submitted   2017-02-23
Author: Lorenz Gerber (`dv15lgr@cs.umu.se lozger03@student.umu.se`)
Instructor: Lars Karlsson / Mikael Ränner

Contents

1 Introduction

This report is part of the mandatory coursework. It describes the solutions for several chosen exercises from the course book [1].

2 4.1 - Generalization of matrix-vector multiplication

If we keep the same scheme of parallelization as mentioned in the book (outer for...loop), generalzation can be implemented rather easy, bascially in the same way as already shown in exercise 1.1:

```
my_first_i = k * m / p + ( k < m mod p ? k : m mod p )
my_last_i = (k + 1) * m / p + (k + 1 < m mod p ? k + 1 : m mod p)
```

It is not useful to parallelize into n as one thread needs to process as this would create a mutex for acces to shared variables.

3 4.2 - Physical data distribution

The source code for this exercise can be found in appendix **??**. Timing:

4 4.8 - Deadlock

a) Both thread wait for the other to release the lock

b) busy waiting with two flag variables would create the same situation as two mutex locks

c) if two semaphores are used, they act more or less in the same way as mutex locks

The main problem in the mentioned program is that it tries to acquire locks from within critical sections, interleaved.

5 4.11 - Linked list troubles

The linked list implementation from section 4.9.2 of the course book was assumed for this exercise. The situations to be described are shown in figure 1.

a) Here, first both threads have to find the element to be removed, 5 and 6. They will first probably first redefine the linked list pointer from their *pred_p* element to point to the next element of the *curr_p* element instead of the *curr_p* itself. Here it could happend that the thread that wants to delete element 8 will never arrive at its destined element as the other thread has already redefined the pointer. In the next step The pointer from *curr_p* will be set to *null* and the actual element free'd. Here again, one problem that could occur is that the thread to delete element 6 will not arrive at his destination because of a *null* pointer. Another obvious problem is that if the thread to delete element 5 will probably already have redifined the pointer from element 2 to element 8, when the other thread will free element 8. Two delete operations will
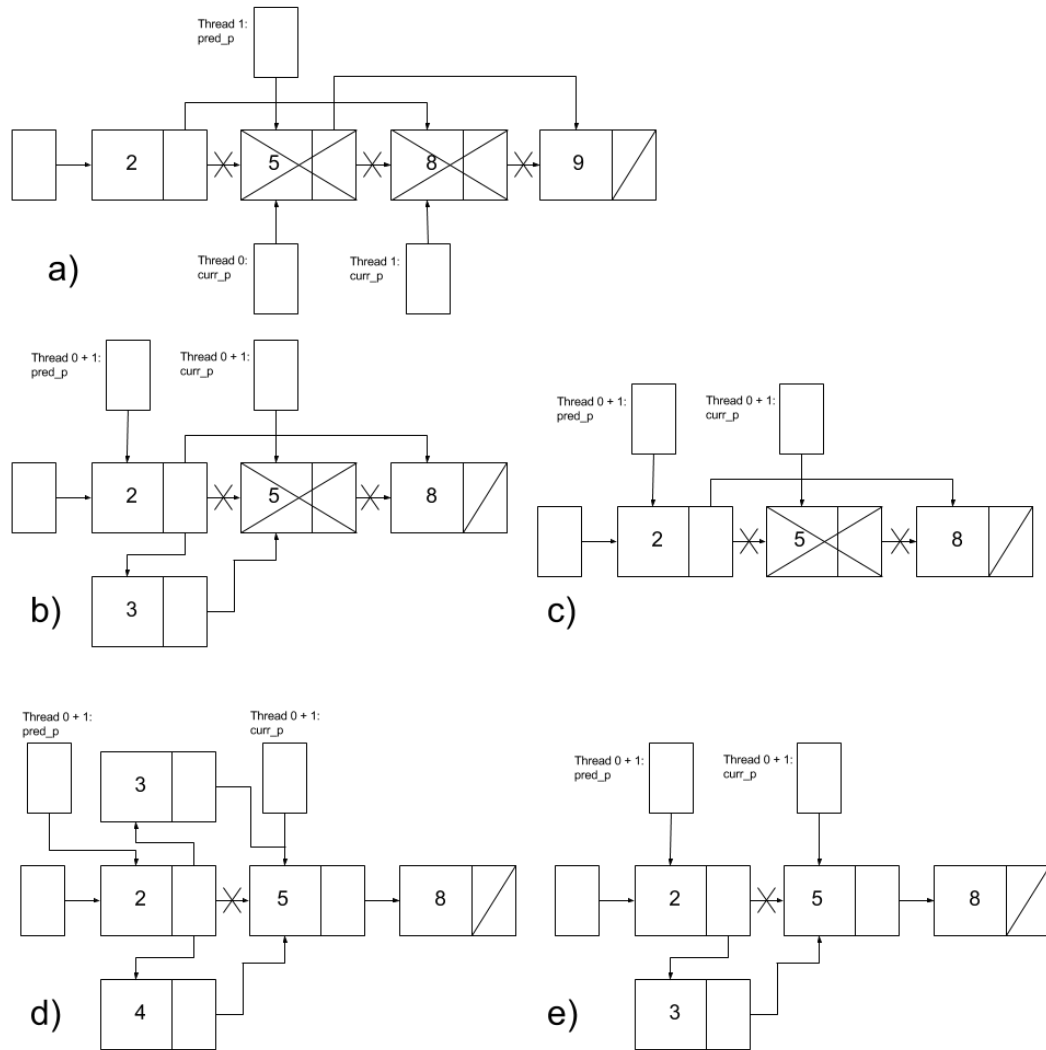
Figure 1: *This graph shows the multithreaded linked list situations from a) to e) that are further described in the text.*

certainly result in either a run time error immediately, or incorrect data structure that probably at a later point in time will lead to a run time error or worse, incorrect results.

b) If an insert and a delete are conducted on the same element, here element 5, the most obvious problem is that the new element 3 will point to a free'd element. Even if element 5 wouldn't be overwritten in the memory yet, it's pointer will be set to *null* anyway, hence the connection to the next list element is lost. Another problem is that both threads will try to redefine the pointer of element 2. Also this situation will probably result in run-time errors such as segmentation fault. If not immediately then probably at a later access. In a rare case when the program will not crash, it seems most likely that the data structure represented will be the linked list 2, 5 with 3 missing.

c) Here one thread wants to delete element 5 while another thread attempts a member operation on element 5. If thread doing a member operation is earlier, it will succeed in his operation, but in respect to the situation after the delete report a wrong value as element 5 does not exist anymore. In some cases, the thread to delete 5 could already have redefined the pointer in element 2 hence the memeber operation would then probably be executed on list element 8 instead of 5. Altough, note, that this could be the intended action. Basically, it will represent the current situation and the data structure is still sane.

d) When two threads try to insert a new element (3 and 4) at the same place (before element 5), it is obvious that the redefination of pointer from element 2 will happen sequential (or else a segmentation fault will happen already here), hence the linked list will incorporate either element 3 or 4. Probably here the chance is big that the program will not crash but that it will produce wrong results.

e) An insert and a member operation from two separate threads on the same list element 5, will most likely result in a wrong result from the member operation in respect to the new data structure. However, in most cases the program should not crash. Depending on timing, it could also happen some times that the member operation will be conducted already at on the new element 3. Note that it is not defined what is correct here, a member function reporting element 3 or 5. This could also be seen as a race condition with undetermined behaviour.

6 4.12 - Linked list insert and delete with read-write lock

No, it's not safe. Another thread could also have the read-lock and request the write lock at the same time. Then only one of two threads will initially get the write-lock and be able to modify the list. When the second thread finally get's the write lock, the list is probably already modified from the first write access and the program will crash.

7 4.17 - False sharing

a) The y vector's length is 8 doubles, hence 64 byes. The minimum number of cache lines is therefore one.

b) It is assumed that the vector is a consecutive memory address area. The 64 bytes long vector can therefore stretch over maximum two cache lines.

c) Stretching over two cache lines, together of length 128 bytes, a 64 bytes long stretch can be divided in 8 bytes steps in 6 different ways.

d) Four threads on two dual core processors results in choosing 2 from 4 permutation ($\frac{n!}{k!(n-k)!}$). The result is 6.

e) Yes, if the vector is split half between two cache lines and if both threads from one processor access the vector positions from either the beginning to the middle or from the middle to the end.

f) The number of such combinations is $6 \times 8 = 48$.

g) There are six such combinations of thread assignments to the cores so that no false sharing occurs.

8 A4.1 - Histogram

The source code can be found in appendix B. Initially, the serial program from chappter 2 was copied and then modified for pthread support. Basically, the whole data structure is kept global. The data access is not critical. Increasing the count values is controlled by a mutex.

9 A4.3 - Trapezoidal rule

The program was redesigned from the serial trapezoidal-rule program contained in chapte 3 of the book. It was decided that only the loop within the *Trap* function shall be parallelized. Hence the first and last calculation: `integral = (f(a) + f(b))/2.0;` and `integral = integral*h` are executed just once by thread 0. The source code is shown in appendix C. As the modifications for busy waiting and semaphore are minimal, only the critical sections with the respective code are shown below

Busy waiting

```
while(flag != my_rank);
integral += local_integral;
flag = (flag+1)%thread_count;
```

Additionally to the above code, the flag variable has to be set as a global.

Semaphore

```
sem_wait(&bin_sem);
integral += local_integral;
sem_post(&bin_sem);
```

Additionally, the header *semaphore.h* has to be loadet as well as the semaphore variable initialised and destroyed after usage.

9.1 Comparison between Mutex, Busy Waiting and Semaphore

In the presented implementation, there are some subtle differences between the three different protections for the critical section. Mutex can probably be seen as the standard way providing access to a critical section. Busy waiting in the current implementation allows access only in a predefined order, along ascending rank. This seems for the present application an acceptable limitation, however in many cases, it would be unsuitable. A disadvantage of busy waiting is that it eventually could break when it is forgotten to dissable compiler optimization. The implementation of the semaphore was here in a very similar way as the mutex. In my case, I often develop on OSX, it is a disadvantage that *semaphore.h* library will not work here. According to the course book, there is another implementation, named semaphores, that will work on OSX.

10 A4.4 - Fork/join overhead

A short program was written that takes a timing for creating and joining threads. The source code can be found in appendix D.

11 A4.5 - Task queue

Here bits and parts from the *pth_ll_one_mut.c* program in the book were taken to construct a task queue that assigns blocks of work to worker threads. In the presented implementation, the work block is a vector of length 10 that is filled with random values that are then interpreted as list operations according to how *pth_ll_one_mut.c* handled it in the original version. A conditional variable is used to assign work to individual worker threads. Which worker thread to assign to is chosen by a call to a random number function. The source code for the implementation can be found in appendix E. The code for rand.c and timer.h is not shown in the appendix but is needed to compile the program.

References

[1] P.S. Pacheco. *An Introduction to Parallel Programming*. Morgan Kaufman, 2011.

A C Source Code for Exercise 4.2

```c
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include "timer.h"

/* Global variables */
int     thread_count;
int     m, n;
int     input_counter = 0;
int     output_counter = 0;
pthread_mutex_t mutex_input;
pthread_mutex_t mutex_output;
```

```c
pthread_cond_t cond_var_input;
pthread_cond_t cond_var_output;

/* Serial functions */
void Usage(char* prog_name);
void Gen_matrix(double A[], int m, int n);
void Read_matrix(char* prompt, double A[], int m, int n);
void Gen_vector(double x[], int n);
void Read_vector(char* prompt, double x[], int n);
void Print_matrix(char* title, double A[], int m, int n);
void Print_vector(char* title, double y[], double m);

/* Parallel function */
void *Pth_mat_vect(void* rank);



/*-------------------------------------------------------------------*/
int main(int argc, char* argv[]) {

   long       thread;
   pthread_t* thread_handles;

   if (argc != 4) Usage(argv[0]);
   thread_count = strtol(argv[1], NULL, 10);
   m = strtol(argv[2], NULL, 10);
   n = strtol(argv[3], NULL, 10);

#  ifdef DEBUG
   printf("thread_count =  %d, m = %d, n = %d\n", thread_count, m, n);
#  endif

   thread_handles = malloc(thread_count*sizeof(pthread_t));

   for (thread = 0; thread < thread_count; thread++)
      pthread_create(&thread_handles[thread], NULL,
         Pth_mat_vect, (void*) thread);

   for (thread = 0; thread < thread_count; thread++)
      pthread_join(thread_handles[thread], NULL);

    return 0;
}  /* main */


void Usage (char* prog_name) {
   fprintf(stderr, "usage: %s <thread_count> <m> <n>\n", prog_name);
   exit(0);
}  /* Usage */
```

```c
void Read_matrix(char* prompt, double A[], int m, int n) {
   int            i, j;

   printf("%s\n", prompt);
   for (i = 0; i < m; i++)
      for (j = 0; j < n; j++)
         scanf("%lf", &A[i*n+j]);
}  /* Read_matrix */


void Gen_matrix(double A[], int m, int n) {
   int i, j;
   for (i = 0; i < m; i++)
      for (j = 0; j < n; j++)
         A[i*n+j] = random()/((double) RAND_MAX);
}  /* Gen_matrix */


void Gen_vector(double x[], int n) {
   int i;
   for (i = 0; i < n; i++)
      x[i] = random()/((double) RAND_MAX);
}

void Read_vector(char* prompt, double x[], int n) {
   int   i;

   printf("%s\n", prompt);
   for (i = 0; i < n; i++)
      scanf("%lf", &x[i]);
}


void *Pth_mat_vect(void* rank) {
  double* A;
  double* x;
  double* y;
  long my_rank = (long) rank;
  int i;
  int j;
  int local_m = m/thread_count;
  int my_first_row = 0; //my_rank*local_m;
  int my_last_row = m/thread_count; //my_first_row + local_m;
  register int sub = my_first_row*n;
  double start, finish;
  double temp;
```

```c
# ifdef DEBUG
  printf("Thread %ld > local_m = %d, sub = %d\n",
        my_rank, local_m, sub);
# endif
  // scheduling data input
  pthread_mutex_lock(&mutex_input);

  while(input_counter != my_rank){
    pthread_cond_wait(&cond_var_input, &mutex_input);
  }

  A = malloc(m/thread_count*n*sizeof(double));
  x = malloc(n*sizeof(double));
  y = malloc(m/thread_count*sizeof(double));

  Gen_matrix(A, local_m, n);
# ifdef DEBUG
  Print_matrix("We generated", A, local_m, n);
# endif

  Gen_vector(x, n);
# ifdef DEBUG
  Print_vector("We generated", x, n);
# endif


  input_counter++;
  pthread_cond_broadcast(&cond_var_input);
  pthread_mutex_unlock(&mutex_input);

  GET_TIME(start);
  for (i = my_first_row; i < my_last_row; i++) {
     y[i] = 0.0;
     for (j = 0; j < n; j++) {
        temp = A[sub++];
        temp *= x[j];
        y[i] += temp;
     }
  }
  GET_TIME(finish);

  pthread_mutex_lock(&mutex_output);

  while(output_counter != my_rank){
    pthread_cond_wait(&cond_var_output, &mutex_output);
  }
  printf("Thread %ld > Elapsed time = %e seconds\n",
     my_rank, finish - start);
```

```
# ifdef DEBUG
  Print_vector("The Product is", y, m/thread_count);
# endif


  output_counter++;
  pthread_cond_broadcast(&cond_var_output);
  pthread_mutex_unlock(&mutex_input);


  free(A);
  free(x);
  free(y);

  return NULL;
}


void Print_matrix( char* title, double A[], int m, int n) {
   int   i, j;

   printf("%s\n", title);
   for (i = 0; i < m; i++) {
      for (j = 0; j < n; j++)
         printf("%6.3f ", A[i*n + j]);
      printf("\n");
   }
}


void Print_vector(char* title, double y[], double m) {
   int   i;

   printf("%s\n", title);
   for (i = 0; i < m; i++)
      printf("%6.3f ", y[i]);
   printf("\n");
}
```

B   C Source Code for Exercise A4.1

```
/* File:       pth_histogram.c
 * Purpose:    Build a histogram from some random data
 *
 * Compile:    gcc -g -Wall -o pth_histogram pth_histogram.c
 * Run:        ./pth_histogram <bin_count> <min_meas> <max_meas>
```

```
 *                  <data_count> <thread_count>
 *
 * Input:     None
 * Output:    A histogram with X's showing the number of measurements
 *            in each bin
 *
 * Notes:
 * 1.  Actual measurements y are in the range min_meas <= y < max_meas
 * 2.  bin_counts[i] stores the number of measurements x in the range
 * 3.  bin_maxes[i-1] <= x < bin_maxes[i] (bin_maxes[-1] = min_meas)
 * 4.  DEBUG compile flag gives verbose output
 * 5.  The program will terminate if either the number of command line
 *     arguments is incorrect or if the search for a bin for a
 *     measurement fails.
 *
 * IPP:  Section 2.7.1 (pp. 66 and ff.)
 */
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

/* Global variables */
int thread_count;
int bin_count;
float min_meas, max_meas;
float* bin_maxes;
int* bin_counts;
int data_count;
float* data;
pthread_mutex_t mutex;

/* Serial functions */
void Usage(char prog_name[]);

void Get_args(
      char*    argv[]        /* in  */,
      int*     bin_count_p   /* out */,
      float*   min_meas_p    /* out */,
      float*   max_meas_p    /* out */,
      int*     data_count_p  /* out */,
      int*     thread_count  /* out */);

void Gen_data(
      float    min_meas     /* in  */,
      float    max_meas     /* in  */,
      float    data[]       /* out */,
      int      data_count   /* in  */);
```

```
void Gen_bins(
      float min_meas       /* in  */,
      float max_meas       /* in  */,
      float bin_maxes[]    /* out */,
      int   bin_counts[]   /* out */,
      int   bin_count      /* in  */);

int Which_bin(
      float    data          /* in */,
      float    bin_maxes[]  /* in */,
      int      bin_count     /* in */,
      float    min_meas      /* in */);

void Print_histo(
      float    bin_maxes[]   /* in */,
      int      bin_counts[]  /* in */,
      int      bin_count     /* in */,
      float    min_meas      /* in */);

/* Parallel function */
void *Pth_count(void* rank);




int main(int argc, char* argv[]) {

   long thread;
   pthread_t* thread_handles;

   /* Check and get command line args */
   if (argc != 6) Usage(argv[0]);
   Get_args(argv, &bin_count, &min_meas, &max_meas, &data_count, &thread_count);

   /* Allocate arrays needed */
   bin_maxes = malloc(bin_count*sizeof(float));
   bin_counts = malloc(bin_count*sizeof(int));
   data = malloc(data_count*sizeof(float));

   /* Generate the data */
   Gen_data(min_meas, max_meas, data, data_count);

   /* Create bins for storing counts */
   Gen_bins(min_meas, max_meas, bin_maxes, bin_counts, bin_count);


   thread_handles = malloc(thread_count*sizeof(pthread_t));

   for (thread = 0; thread < thread_count; thread++)
```

```c
        pthread_create(&thread_handles[thread], NULL,
         Pth_count, (void*) thread);

     for (thread = 0; thread < thread_count; thread++)
        pthread_join(thread_handles[thread], NULL);




#  ifdef DEBUG
     printf("bin_counts = ");
     for (i = 0; i < bin_count; i++)
        printf("%d ", bin_counts[i]);
     printf("\n");
#  endif

     /* Print the histogram */
     Print_histo(bin_maxes, bin_counts, bin_count, min_meas);

     free(data);
     free(bin_maxes);
     free(bin_counts);
     return 0;

}  /* main */

void *Pth_count(void* rank){
  /* which data to access */
  int i, bin;
  long my_rank = (long) rank;
  int first_i = data_count/thread_count* (int) my_rank;
  int last_i = data_count/thread_count* ((int) my_rank + 1)-1;

  /* Count number of values in each bin */
  for (i = first_i; i < last_i; i++) {
    bin = Which_bin(data[i], bin_maxes, bin_count, min_meas);
    pthread_mutex_lock(&mutex);
    bin_counts[bin]++;
    pthread_mutex_unlock(&mutex);
  }

  return NULL;

}
```

```c
/*-------------------------------------------------------------------
 * Function:  Usage
 * Purpose:   Print a message showing how to run program and quit
 * In arg:    prog_name:  the name of the program from the command line
 */
void Usage(char prog_name[] /* in */) {
   fprintf(stderr, "usage: %s ", prog_name);
   fprintf(stderr, "<bin_count> <min_meas> <max_meas> <data_count> <thread_count>\n");
   exit(0);
}  /* Usage */


/*-------------------------------------------------------------------
 * Function:  Get_args
 * Purpose:   Get the command line arguments
 * In arg:    argv:  strings from command line
 * Out args:  bin_count_p:   number of bins
 *            min_meas_p:    minimum measurement
 *            max_meas_p:    maximum measurement
 *            data_count_p:  number of measurements
 */
void Get_args(
      char*    argv[]        /* in  */,
      int*     bin_count_p   /* out */,
      float*   min_meas_p    /* out */,
      float*   max_meas_p    /* out */,
      int*     data_count_p  /* out */,
      int*     thread_count  /* out */) {

   *bin_count_p = strtol(argv[1], NULL, 10);
   *min_meas_p = strtof(argv[2], NULL);
   *max_meas_p = strtof(argv[3], NULL);
   *data_count_p = strtol(argv[4], NULL, 10);
   *thread_count = strtol(argv[5], NULL, 10);

#  ifdef DEBUG
   printf("bin_count = %d\n", *bin_count_p);
   printf("min_meas = %f, max_meas = %f\n", *min_meas_p, *max_meas_p);
   printf("data_count = %d\n", *data_count_p);
   printf("thread_count = %d\n", *thread_count);
#  endif
}  /* Get_args */


/*-------------------------------------------------------------------
 * Function:  Gen_data
 * Purpose:   Generate random floats in the range min_meas <= x < max_meas
 * In args:   min_meas:    the minimum possible value for the data
```

```
 *             max_meas:   the maximum possible value for the data
 *             data_count: the number of measurements
 * Out arg:   data:       the actual measurements
 */
void Gen_data(
        float   min_meas    /* in  */,
        float   max_meas    /* in  */,
        float   data[]      /* out */,
        int     data_count  /* in  */) {
   int i;

   srandom(0);
   for (i = 0; i < data_count; i++)
      data[i] = min_meas + (max_meas - min_meas)*random()/((double) RAND_MAX);

#  ifdef DEBUG
   printf("data = ");
   for (i = 0; i < data_count; i++)
      printf("%4.3f ", data[i]);
   printf("\n");
#  endif
}  /* Gen_data */


/*-------------------------------------------------------------------
 * Function:  Gen_bins
 * Purpose:   Compute max value for each bin, and store 0 as the
 *            number of values in each bin
 * In args:   min_meas:   the minimum possible measurement
 *            max_meas:   the maximum possible measurement
 *            bin_count:  the number of bins
 * Out args:  bin_maxes:  the maximum possible value for each bin
 *            bin_counts: the number of data values in each bin
 */
void Gen_bins(
     float min_meas      /* in  */,
     float max_meas      /* in  */,
     float bin_maxes[]   /* out */,
     int   bin_counts[]  /* out */,
     int   bin_count     /* in  */) {
   float bin_width;
   int   i;

   bin_width = (max_meas - min_meas)/bin_count;

   for (i = 0; i < bin_count; i++) {
      bin_maxes[i] = min_meas + (i+1)*bin_width;
      bin_counts[i] = 0;
```

```
   }

#  ifdef DEBUG
   printf("bin_maxes = ");
   for (i = 0; i < bin_count; i++)
      printf("%4.3f ", bin_maxes[i]);
   printf("\n");
#  endif
}  /* Gen_bins */


/*-------------------------------------------------------------------
 * Function:  Which_bin
 * Purpose:   Use binary search to determine which bin a measurement
 *            belongs to
 * In args:   data:       the current measurement
 *            bin_maxes:  list of max bin values
 *            bin_count:  number of bins
 *            min_meas:   the minimum possible measurement
 * Return:    the number of the bin to which data belongs
 * Notes:
 * 1.  The bin to which data belongs satisfies
 *
 *            bin_maxes[i-1] <= data < bin_maxes[i]
 *
 *     where, bin_maxes[-1] = min_meas
 * 2.  If the search fails, the function prints a message and exits
 */
int Which_bin(
      float   data          /* in */,
      float   bin_maxes[]   /* in */,
      int     bin_count     /* in */,
      float   min_meas      /* in */) {
   int bottom = 0, top =  bin_count-1;
   int mid;
   float bin_max, bin_min;

   while (bottom <= top) {
      mid = (bottom + top)/2;
      bin_max = bin_maxes[mid];
      bin_min = (mid == 0) ? min_meas: bin_maxes[mid-1];
      if (data >= bin_max)
         bottom = mid+1;
      else if (data < bin_min)
         top = mid-1;
      else
         return mid;
   }
```

```
   /* Whoops! */
   fprintf(stderr, "Data = %f doesn't belong to a bin!\n", data);
   fprintf(stderr, "Quitting\n");
   exit(-1);
}  /* Which_bin */


/*-------------------------------------------------------------------
 * Function:  Print_histo
 * Purpose:   Print a histogram.  The number of elements in each
 *            bin is shown by an array of X's.
 * In args:   bin_maxes:   the max value for each bin
 *            bin_counts:  the number of elements in each bin
 *            bin_count:   the number of bins
 *            min_meas:    the minimum possible measurment
 */
void Print_histo(
        float  bin_maxes[]   /* in */,
        int    bin_counts[]  /* in */,
        int    bin_count     /* in */,
        float  min_meas      /* in */) {
   int i, j;
   float bin_max, bin_min;

   for (i = 0; i < bin_count; i++) {
      bin_max = bin_maxes[i];
      bin_min = (i == 0) ? min_meas: bin_maxes[i-1];
      printf("%.3f-%.3f:\t", bin_min, bin_max);
      for (j = 0; j < bin_counts[i]; j++)
         printf("X");
      printf("\n");
   }
}  /* Print_histo */
```

C   C Source Code for Exercise A4.3

```
/* File:    trap.c
 * Purpose: Calculate definite integral using trapezoidal
 *          rule.
 *
 * Input:   a, b, n, thread_count
 * Output:  Estimate of integral from a to b of f(x)
 *          using n trapezoids.
 *
 * Compile: gcc -g -Wall -o trap trap.c
```

```
 * Usage:    ./trap
 *
 * Note:     The function f(x) is hardwired. thread_count should
 *           the current version be integer divisible by n-1
 *
 * IPP:      Section 3.2.1 (pp. 94 and ff.) and 5.2 (p. 216)
 */

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

/* global variables */
int thread_count;
double  integral;   /* Store result in integral   */
double  a, b;       /* Left and right endpoints   */
int     n;          /* Number of trapezoids       */
double  h;          /* Height of trapezoids       */
pthread_mutex_t mutex;

/* serial functions */
double f(double x);     /* Function we're integrating */
void Trap(double a, double b, int n, double h);

/* parallel function */
void *Pth_interpol(void* rank);


int main(void) {

   printf("Enter a, b, n and thread_count\n");
   scanf("%lf", &a);
   scanf("%lf", &b);
   scanf("%d", &n);
   scanf("%d", &thread_count);

   h = (b-a)/n;
   Trap(a, b, n, h);

   printf("With n = %d trapezoids, our estimate\n", n);
   printf("of the integral from %f to %f = %.15f\n",
       a, b, integral);

   return 0;
}  /* main */

/*------------------------------------------------------------
 * Function:    Trap
```

```
 * Purpose:     Estimate integral from a to b of f using trap rule and
 *              n trapezoids
 * Input args:  a, b, n, h
 * Return val:  Estimate of the integral
 */
void Trap(double a, double b, int n, double h) {
   long thread;
   pthread_t* thread_handles;

   integral = (f(a) + f(b))/2.0;

   /* here comes the parallel part */

   thread_handles = malloc(thread_count*sizeof(pthread_t));

   for (thread = 0; thread < thread_count; thread++)
     pthread_create(&thread_handles[thread], NULL,
    Pth_interpol, (void*) thread);

   for (thread = 0; thread < thread_count; thread++)
     pthread_join(thread_handles[thread], NULL);

   integral = integral*h;

   }  /* Trap */

void *Pth_interpol(void* rank){
  /* not the first and not the last */
  long my_rank = (long) rank;
  int first_k = (n-1) / thread_count * my_rank + 1 ;
  int last_k= (n-1) / thread_count * (my_rank+1);
  double local_integral = 0;
  int k;

  //printf("first %d last %d\n", first_k, last_k);

  for (k = first_k; k <= last_k; k++) {
     local_integral += f(a+k*h);
  }

  /* updating global sum */
  pthread_mutex_lock(&mutex);
  integral += local_integral;
  pthread_mutex_unlock(&mutex);

  return NULL;

}
```

```
/*----------------------------------------------------------------
 * Function:    f
 * Purpose:     Compute value of function to be integrated
 * Input args:  x
 */
double f(double x) {
   double return_val;

   return_val = x*x;
   return return_val;
}  /* f */
```

D   C Source Code for Exercise A4.4

```c
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include "timer.h"

/* Global variables */
int      thread_count = 1;
void *Pth_empty(void* rank);



int main(int argc, char* argv[]) {

  /* local variables */
  long       thread;
  pthread_t* thread_handles;
  double start, finish;

  /* outer loop to test for increasing number of threads */
  for(int j = thread_count; j <= 8; j=j+2){

    /* inner loop to repeat 100 measurments */
    for (int i = 0; i < 100; i++){

      GET_TIME(start);

      /* assign, create and join j threads */
      thread_handles = malloc(j*sizeof(pthread_t));

      for (thread = 0; thread < j; thread++){
        pthread_create(&thread_handles[thread], NULL,
    Pth_empty, (void*) thread);
```

```
        }

        for (thread = 0; thread < j; thread++){
          pthread_join(thread_handles[thread], NULL);
        }

        GET_TIME(finish);
        free(thread_handles);

        /* output results */
        printf("%d %e\n", j, (finish - start)/j);
      }
   }

   return 0;
}

/* dummy/ empty thread function */
void *Pth_empty(void* rank){
   return NULL;
}
```

E C Source Code for Exercise A4.5

```
/* File:     pth_thread_pool.c
 *
 * Purpose:  Implement a thread pool task queue with sorted linked list
             operations.
 *
 * Compile:  gcc -g -Wall -o pth_thread_pool pth_thread_pool.c
 *               my_rand.c -lpthread
 *           needs timer.h and my_rand.h
 *
 * Usage:    ./pth_thread_cound <thread_count>
 * Input:    total number of keys inserted by main thread
 *           total number of work blocks to assign
 *           percent of ops that are searches and inserts (remaining ops
 *             are deletes.
 * Output:   Elapsed time to carry out the ops
 *
 */
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include "my_rand.h"
#include "timer.h"
```

```c
/* Random ints are less than MAX_KEY */
const int MAX_KEY = 100000000;

/* Struct for list nodes */
struct list_node_s {
   int     data;
   struct list_node_s* next;
};

/* Shared variables */
struct     list_node_s* head = NULL;
int        thread_count;
int        total_ops;
double     insert_percent;
double     search_percent;
double     delete_percent;
pthread_mutex_t mutex;
pthread_mutex_t count_mutex;
pthread_mutex_t cond_mutex;
pthread_cond_t cond_var;
int        member_total=0, insert_total=0, delete_total=0;
double work_block[10] = {0};
int no_work_blocks;
int finish_job = 0;
int thread_assignment = 0;

/* Setup and cleanup */
void       Usage(char* prog_name);
void       Get_input(int* inserts_in_main_p);

/* Thread function */
void*      Thread_work(void* rank);

/* List operations */
int        Insert(int value);
void       Print(void);
int        Member(int value);
int        Delete(int value);
void       Free_list(void);
int        Is_empty(void);

/*-------------------------------------------------------------------*/
int main(int argc, char* argv[]) {
   long i;
   int key, success, attempts;
   pthread_t* thread_handles;
   int inserts_in_main;
```

```
unsigned seed = 1;
double start, finish;

int work_counter;
int worker_id = 0;

/* take care of command line args */
if (argc != 2) Usage(argv[0]);
thread_count = strtol(argv[1],NULL,10);

/* Get user input */
Get_input(&inserts_in_main);

/* Try to insert inserts_in_main keys, but give up after */
/* 2*inserts_in_main attempts.                           */
i = attempts = 0;
while ( i < inserts_in_main && attempts < 2*inserts_in_main ) {
   key = my_rand(&seed) % MAX_KEY;
   success = Insert(key);
   attempts++;
   if (success) i++;
}
printf("Inserted %ld keys in empty list\n", i);


/* Prepare threads and mutexes */
thread_handles = malloc((thread_count-1)*sizeof(pthread_t));
pthread_mutex_init(&mutex, NULL);
pthread_mutex_init(&count_mutex, NULL);

GET_TIME(start);
/* start threads */
for (i = 1; i <= thread_count; i++)
   pthread_create(&thread_handles[i], NULL, Thread_work, (void*) i);

/* Generate work blocks */
for (work_counter = 0; work_counter < no_work_blocks; work_counter++){

  /* check conditional, if zero, proceed */
  pthread_mutex_lock(&cond_mutex);
  while(thread_assignment != 0){
    pthread_cond_wait(&cond_var, &cond_mutex);
  }

  /* fill work block */
  for(i = 0; i < 10; i++){
    work_block[i] = my_drand(&seed);
  }
```

```
      /* assign thread for work */
      worker_id = (my_rand(&seed) % (thread_count-1)) + 1;
      thread_assignment = worker_id;
      pthread_mutex_unlock(&cond_mutex);

   }

   /* wait until last thread is in work loop */
   pthread_mutex_lock(&cond_mutex);
   while(thread_assignment != 0){
     pthread_cond_wait(&cond_var, &cond_mutex);
   }

   /* set finish_job to true */
   finish_job = 1;

   /* broadcast threads */
   pthread_cond_broadcast(&cond_var);
   pthread_mutex_unlock(&cond_mutex);

   /* join threads after done work */
   for (i = 1; i <= thread_count; i++){
     pthread_join(thread_handles[i], NULL);
   }

   /* report results */
   GET_TIME(finish);
   printf("Elapsed time = %e seconds\n", finish - start);
   printf("Total ops = %d\n", no_work_blocks * 10);
   printf("member ops = %d\n", member_total);
   printf("insert ops = %d\n", insert_total);
   printf("delete ops = %d\n", delete_total);

   /* clean up */
   Free_list();
   pthread_mutex_destroy(&mutex);
   pthread_mutex_destroy(&count_mutex);
   pthread_mutex_destroy(&cond_mutex);
   pthread_cond_destroy(&cond_var);
   free(thread_handles);

   return 0;
}  /* main */


/*-------------------------------------------------------------------*/
void Usage(char* prog_name) {
```

```
      fprintf(stderr, "usage: %s <thread_count>\n", prog_name);
      exit(0);
}  /* Usage */

/*-------------------------------------------------------------------*/
void Get_input(int* inserts_in_main_p) {

   printf("How many keys should be inserted in the main thread?\n");
   scanf("%d", inserts_in_main_p);
   printf("How many work blocks of each 10 operations should be executed?\n");
   scanf("%d", &no_work_blocks);
   printf("Percent of ops that should be searches? (between 0 and 1)\n");
   scanf("%lf", &search_percent);
   printf("Percent of ops that should be inserts? (between 0 and 1)\n");
   scanf("%lf", &insert_percent);
   delete_percent = 1.0 - (search_percent + insert_percent);
}  /* Get_input */

/*-------------------------------------------------------------------*/
/* Insert value in correct numerical location into list */
/* If value is not in list, return 1, else return 0 */
int Insert(int value) {
   struct list_node_s* curr = head;
   struct list_node_s* pred = NULL;
   struct list_node_s* temp;
   int rv = 1;

   while (curr != NULL && curr->data < value) {
      pred = curr;
      curr = curr->next;
   }

   if (curr == NULL || curr->data > value) {
      temp = malloc(sizeof(struct list_node_s));
      temp->data = value;
      temp->next = curr;
      if (pred == NULL)
         head = temp;
      else
         pred->next = temp;
   } else { /* value in list */
      rv = 0;
   }

   return rv;
}  /* Insert */

/*-------------------------------------------------------------------*/
```

```
void Print(void) {
   struct list_node_s* temp;

   printf("list = ");

   temp = head;
   while (temp != (struct list_node_s*) NULL) {
      printf("%d ", temp->data);
      temp = temp->next;
   }
   printf("\n");
}  /* Print */


/*----------------------------------------------------------------*/
int  Member(int value) {
   struct list_node_s* temp;

   temp = head;
   while (temp != NULL && temp->data < value)
      temp = temp->next;

   if (temp == NULL || temp->data > value) {
#     ifdef DEBUG
      printf("%d is not in the list\n", value);
#     endif
      return 0;
   } else {
#     ifdef DEBUG
      printf("%d is in the list\n", value);
#     endif
      return 1;
   }
}  /* Member */

/*----------------------------------------------------------------*/
/* Deletes value from list */
/* If value is in list, return 1, else return 0 */
int Delete(int value) {
   struct list_node_s* curr = head;
   struct list_node_s* pred = NULL;
   int rv = 1;

   /* Find value */
   while (curr != NULL && curr->data < value) {
      pred = curr;
      curr = curr->next;
   }
```

```c
      if (curr != NULL && curr->data == value) {
         if (pred == NULL) { /* first element in list */
            head = curr->next;
#        ifdef DEBUG
            printf("Freeing %d\n", value);
#        endif
            free(curr);
         } else {
            pred->next = curr->next;
#        ifdef DEBUG
            printf("Freeing %d\n", value);
#        endif
            free(curr);
         }
      } else { /* Not in list */
         rv = 0;
      }

      return rv;
}  /* Delete */

/*-------------------------------------------------------------------*/
void Free_list(void) {
   struct list_node_s* current;
   struct list_node_s* following;

   if (Is_empty()) return;
   current = head;
   following = current->next;
   while (following != NULL) {
#     ifdef DEBUG
      printf("Freeing %d\n", current->data);
#     endif
      free(current);
      current = following;
      following = current->next;
   }
#  ifdef DEBUG
   printf("Freeing %d\n", current->data);
#  endif
   free(current);
}  /* Free_list */

/*-------------------------------------------------------------------*/
int  Is_empty(void) {
   if (head == NULL)
      return 1;
```

```
       else
          return 0;
}  /* Is_empty */

/*-------------------------------------------------------------------*/
void* Thread_work(void* rank) {
   long my_rank = (long) rank;
   int i, val;
   unsigned seed = my_rank + 1;
   int my_member=0, my_insert=0, my_delete=0;
   double local_work_block[10] = {0};

   while(1){

     /* go to conditional sleep */
     pthread_mutex_lock(&cond_mutex);

     while(thread_assignment != my_rank && finish_job == 0){
       pthread_cond_wait(&cond_var, &cond_mutex);
     }

     /* check project finished */
     if (finish_job){
       return NULL;
     }

     /* if my turn read, global work block to local work block */
     for (i = 0; i < 10; i++){
       local_work_block[i] = work_block[i];
     }

     /* set thread_assignment back to zero */
     thread_assignment = 0;

     /* do work */
     for (i = 0; i < 10; i++) {
       val = my_rand(&seed) % MAX_KEY;
       if (local_work_block[i] < search_percent) {
           pthread_mutex_lock(&mutex);
           Member(val);
           pthread_mutex_unlock(&mutex);
           my_member++;
         } else if (local_work_block[i] < search_percent + insert_percent) {
           pthread_mutex_lock(&mutex);
           Insert(val);
           pthread_mutex_unlock(&mutex);
           my_insert++;
         } else { /* delete */
```

```
            pthread_mutex_lock(&mutex);
            Delete(val);
            pthread_mutex_unlock(&mutex);
            my_delete++;
         }

      }  /* for */

      /* book keeping for results */
      pthread_mutex_lock(&count_mutex);
      member_total += my_member;
      insert_total += my_insert;
      delete_total += my_delete;

      pthread_mutex_unlock(&count_mutex);
      my_member = my_insert = my_delete = 0;

   } /* while */


   return NULL;
}  /* Thread_work */
```